

# EC708 Discussion 8

## Numerical Optimization

Yan Liu<sup>1</sup>

Department of Economics  
Boston University

March 18, 2022

---

<sup>1</sup>Parts of the materials are from Train (2009), Kochenderfer and Wheeler (2019), and teaching slides of Shuowen Chen, Jean-Jacques Forneron, and Ryan Tibshirani.

# Outline

## 1 Optimization Problem

## 2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

## 3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and  
Davidon-Fletcher-Powell (DFP)

## 4 Stochastic Gradient Descent

## 5 Comparison Methods: Nelder-Mead Algorithm

# Optimization Problem

## MLE

Consider maximizing the log-likelihood function


$$\bar{\ell}_T(\theta) = \sum_{t=1}^T \ell_t(\theta),$$

where  $\theta \in \Theta \subset \mathbb{R}^K$ .

- **Goal:** find  $\hat{\theta}_T = \arg \max_{\theta \in \Theta} \bar{\ell}_T(\theta)$ .
- To utilize minimization packages in practice, we usually work with  $-\bar{\ell}_T(\theta)$ .

# Optimization Problem

## Graphical Illustration of MLE



Finding  $\hat{\theta}_T$  is a hill-climbing process:

- 1 Specify a starting value  $\theta_0$ .
- 2 Each step  $k$  moves to a new value  $\theta_{k+1}$  at which  $\bar{\ell}_T(\theta)$  is higher than at the current value  $\theta_k$ .
- 3 Keep climbing until no further increase can be found.

# Optimization Problem

## Gradient and Hessian

Gradient in step  $k$ :

$$g_k = \left( \frac{\partial \bar{\ell}_T(\theta)}{\partial \theta} \right)_{\theta_k}.$$

Hessian in step  $k$ :

$$H_k = \left( \frac{\partial g_k}{\partial \theta'} \right)_{\theta_k} = \left( \frac{\partial^2 \bar{\ell}_T(\theta)}{\partial \theta \partial \theta'} \right)_{\theta_k}.$$

The gradient tells use **in what direction** to climb, and the Hessian can help us to know **how far** to climb.

# Optimization Problem

## Numerical Evaluation of Derivatives

Recall the definition of first-order partial derivative of function  $f$  at  $\theta$ :

$$\frac{\partial f(\theta)}{\partial \theta_j} = \lim_{h \rightarrow 0} \frac{f(\theta_1, \dots, \theta_j + h, \dots, \theta_K) - f(\theta_1, \dots, \theta_j, \dots, \theta_K)}{h}$$

for  $j = 1, \dots, K$ . Numerically, we can approximate it by calculating

$$f_j(\theta) = \frac{f(\theta + he_j) - f(\theta)}{h},$$

where  $e_j = (0, \dots, 0, 1, 0, \dots, 0)$  is the unit vector with 1 in position  $j$  and  $h$  is the step size. In practice, a more accurate approximation is

$$f_j(\theta) = \frac{f(\theta + he_j) - f(\theta - he_j)}{2h}.$$

# Outline

## 1 Optimization Problem

## 2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

## 3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and  
Davidon-Fletcher-Powell (DFP)

## 4 Stochastic Gradient Descent

## 5 Comparison Methods: Nelder-Mead Algorithm

# Newton-Raphson Method

## Algorithm

Take a second-order Taylor's approximation of  $\bar{\ell}_T(\theta_{k+1})$  around  $\bar{\ell}_T(\theta_k)$ :

$$\bar{\ell}_T(\theta_{k+1}) = \bar{\ell}_T(\theta_k) + (\theta_{k+1} - \theta_k)' g_k + \frac{1}{2} (\theta_{k+1} - \theta_k)' H_k (\theta_{k+1} - \theta_k).$$

Find  $\theta_{k+1}$  that maximizes this approximation:


$$g_k + H_k(\theta_{k+1} - \theta_k) = 0 \Rightarrow \theta_{k+1} = \theta_k - \underbrace{H_k^{-1}}_{\text{step size}} \cdot \underbrace{g_k}_{\text{direction}}.$$

- Iterate until convergence, which can be defined in many ways:
  - $\bar{\ell}_T(\theta_{k+1})$  close to  $\bar{\ell}_T(\theta_k)$
  - $\theta_{k+1}$  close to  $\theta_k$
  - $g_{k+1}$  close to  $g_k$
- If  $\bar{\ell}_T(\theta)$  were **exactly quadratic**, then Newton-Raphson reaches the maximum in **one step** from any starting value.



# Newton-Raphson Method

## Step Size




- It is possible for Newton-Raphson to step past the maximum and move to a lower  $\bar{\ell}_T(\theta)$ .
- To ensure each step provides an increase in  $\bar{\ell}_T(\theta)$ , we introduce a scalar step size  $\lambda_k$ :

$$\theta_{k+1} = \theta_k - \lambda_k H_k^{-1} g_k.$$

# Newton-Raphson Method

## Determining the Step Size: Backtracking Line Search



Perform step-size adjustment in each step. Start with  $\lambda_k = 1$ .

- If  $\bar{\ell}_T(\theta_{k+1}) < \bar{\ell}_T(\theta_k)$ , continue halving  $\lambda_k$  until  $\bar{\ell}_T(\theta_{k+1}) > \bar{\ell}_T(\theta_k)$ .
  - A tiny  $\lambda_k$  is a signal that a different iteration procedure is needed.
- If  $\bar{\ell}_T(\theta_{k+1}) > \bar{\ell}_T(\theta_k)$ , continue doubling  $\lambda_k$  as long as doing so further raises  $\bar{\ell}_T(\theta_{k+1})$ .
  - Raising  $\lambda_k$  reduces # of iterations needed to reach the maximum.

# Newton-Raphson Method

## Drawbacks

- Calculations of the Hessian is usually computation-intensive.
  - $\frac{K(K+1)}{2}$  functions to evaluate in each step
  - Numerically calculated Hessian might be ill-behaved (singular).
- Does not guarantee an increase in each step if the log-likelihood function is not globally concave.
  - Hessian may not be negative definite.
  - Remedy: **regularization**. Instead of using  $H_k^{-1}$  directly, use

$$(H_k + \mu_k I_K)^{-1}$$

where  $\mu_k < 0$  guarantees negative definiteness.

# Gauss-Newton Method

Consider the following general nonlinear model:

$$y_t = f(x_t; \theta) + u_t, \quad t = 1, \dots, T$$

- $u_t \sim \text{i.i.d. } N(0, \sigma^2)$ . Assume  $\sigma^2$  is known here.
- $x_t$ :  $M \times 1$  exogenous regressors.
- $f(\cdot)$ : some function satisfying some regularity conditions.

# Gauss-Newton Method

Due to the normality assumption, we have the log-likelihood function

$$\bar{\ell}_T(\theta) = -\frac{T}{2} \ln 2\pi - \frac{T}{2} \ln \sigma^2 - \underbrace{\sum_{t=1}^T (y_t - f(x_t; \theta))^2}_{\equiv S(\theta)}.$$

It suffices to work with  $S(\theta)$ . Its first and second derivatives w.r.t.  $\theta$  are

$$g(\theta) = 2 \sum_{t=1}^T \frac{\partial u_t}{\partial \theta} u_t, \quad H(\theta) = 2 \sum_{t=1}^T \left[ \frac{\partial u_t}{\partial \theta} \frac{\partial u_t}{\partial \theta'} + \frac{\partial^2 u_t}{\partial \theta \partial \theta'} u_t \right].$$

In  $H(\theta)$ , the blue term is usually small relative to the red term, so we **neglect** it and use the following:

$$\theta_{k+1} = \theta_k - \left[ \sum_{t=1}^T \frac{\partial u_t}{\partial \theta} \frac{\partial u_t}{\partial \theta'} \right]^{-1} \bigg|_{\theta=\theta_k} \sum_{t=1}^T \frac{\partial u_t}{\partial \theta} u_t \bigg|_{\theta=\theta_k}.$$

# Gauss-Newton Method

## Remarks:

- This method doesn't compute second-order derivatives.
- Has an OLS interpretation. Let  $z_t = -\partial u_t / \partial \theta$ , then

$$\theta_{k+1} = \theta_k + \left( \sum_{t=1}^T z_t z_t' \right)^{-1} \left| \sum_{t=1}^T z_t u_t \right|_{\theta=\theta_k}.$$

- Similar regularization can be incorporated as in Newton-Raphson:  
Marquart quadratic hill climbing

$$\theta_{k+1} = \theta_k + \left( \sum_{t=1}^T z_t z_t' + \mu I_K \right)^{-1} \left| \sum_{t=1}^T z_t u_t \right|_{\theta=\theta_k}.$$

# Outline

## 1 Optimization Problem

## 2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

## 3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and Davidon-Fletcher-Powell (DFP)

## 4 Stochastic Gradient Descent

## 5 Comparison Methods: Nelder-Mead Algorithm

# Quasi-Newton Methods

- Both main steps in Newton-Raphson method could be expensive:
  - Compute Hessian  $H_k$
  - Solve the system  $H_k \Delta = -g_k$
- Quasi-Newton methods repeat updates of the form

$$\theta_{k+1} = \theta_k - \lambda_k G_k^{-1} g_k$$

for some approximation  $G_k$  of  $H_k$ . We want  $G_k$  to be easy to compute and linear system  $G_k \Delta = -g_k$  to be easy to solve.



# Quasi-Newton Methods

## Berndt-Hall-Hall-Hausman (BHHH)

Berndt et al. (1974) utilize the fact that the objective function is the sum of log likelihoods and propose to use scores to approximate Hessian.

- Score of observation  $t$ :

$$s_t(\theta_k) = \left. \frac{\partial \ell_t(\theta)}{\partial \theta} \right|_{\theta=\theta_k}.$$

- Gradient is the sum of scores:

$$g_k = \sum_{t=1}^T s_t(\theta_k).$$

# Quasi-Newton Methods

## Berndt-Hall-Hall-Hausman (BHHH)

- Outer product of observation  $t$ 's score is the  $K \times K$  matrix

$$s_t(\theta_k)s_t(\theta_k)' = \begin{pmatrix} s_t^1 s_t^1 & s_t^1 s_t^2 & \cdots & s_t^1 s_t^K \\ s_t^2 s_t^1 & s_t^2 s_t^2 & \cdots & s_t^2 s_t^K \\ \vdots & \vdots & & \vdots \\ s_t^K s_t^1 & s_t^K s_t^2 & \cdots & s_t^K s_t^K \end{pmatrix},$$

where  $s_t^j$  is the  $j$ -th element of  $s_t(\theta_k)$ .

- Sum of outer product of scores:

$$G_k = \sum_{t=1}^T s_t(\theta_k)s_t(\theta_k)'.$$

Berndt-Hall-Hall-Hausman (BHHH) update uses  $G_k$  in place of  $-H_k$ :

$$\theta_{k+1} = \theta_k + \lambda_k G_k^{-1} g_k.$$

# Quasi-Newton Methods

Berndt-Hall-Hall-Hausman (BHHH)

Why does BHHH work?

- At maximum,  $G_k$  is the sample variance of scores and thus provides a measure of the log-likelihood functions' curvature, similar to  $H_k$ .
- These ideas are formalized in the **information matrix equality**.
- $G_k$  is far faster to calculate than  $H_k$  and necessarily positive definite.

**Drawbacks:** BHHH can give small steps when far from the maximum because  $G_k$  is not a good approximation to  $-H_k$ .

# Quasi-Newton Methods

## BFGS and DFP

- BHHH uses only information at  $\theta_k$  to determine each step.
- As  $G_k$  already contains information about the Hessian, BFGS and DFP use suitable matrix update to form  $G_{k+1}$ .
- General procedure: In each iteration  $k$ ,
  - 1 Compute Quasi-Newton direction  $\Delta_k = -G_k^{-1}g_k$
  - 2 Determine stepsize  $\lambda_k$  (by backtracking line search)
  - 3 Update  $\theta_{k+1} = \theta_k + \lambda_k \Delta_k$
  - 4 Compute  $G_{k+1}$  from  $G_k$

# Quasi-Newton Methods

BFGS and DFP

Reasonable requirement for  $G_{k+1}$  (motivated by secant method):

$$g_{k+1} = g_k + G_{k+1} \Delta_k.$$

In addition, we want:

- $G_{k+1}$  to be symmetric
- $G_{k+1}$  to preserve positive definiteness (BFGS and DFP deal with convex optimization)

# Quasi-Newton Methods

## Broyden-Fletcher-Goldfarb-Shanno (BFGS)

### **Broyden-Fletcher-Goldfarb-Shanno (BFGS) update:**

BFGS uses a rank-two update of the form:

$$G_{k+1} = G_k + auu' + bvv'.$$

Let  $\gamma_k = g_{k+1} - g_k$ . The secant equation yields

$$\gamma_k - G_k \Delta_k = (au' \Delta_k)u + (bv' \Delta_k)v.$$

Putting  $u = \gamma_k$ ,  $v = G_k \Delta_k$ , and solving for  $a, b$  we get

$$G_{k+1} = G_k - \frac{G_k \Delta_k \Delta_k' G_k}{\Delta_k' G_k \Delta_k} + \frac{\gamma_k \gamma_k'}{\gamma_k' \Delta_k}.$$

- BFGS update is quite cheap:  $O(K^2)$  operations
- BFGS is the algorithm behind Matlab's **fminunc**.

# Quasi-Newton Methods

## Davidon-Fletcher-Powell (DFP)

### Davidon-Fletcher-Powell (DFP) update:

DFP pursues the same idea to update  $G_{k+1}^{-1}$ :

$$G_{k+1}^{-1} = G_k^{-1} + auu' + bvv'.$$

The secant equation yields

$$\Delta_k - G_k^{-1}\gamma_k = (au'\gamma_k)u + (bv'\gamma_k)v.$$

Putting  $u = \Delta_k$ ,  $v = G_k^{-1}\Delta_k$ , and solving for  $a, b$  we get

$$G_{k+1}^{-1} = G_k^{-1} - \frac{G_k^{-1}\gamma_k\gamma_k'G_k^{-1}}{\gamma_k'G_k^{-1}\gamma_k} + \frac{\Delta_k\Delta_k'}{\Delta_k'\gamma_k}.$$

- The role of  $\gamma_k$  and  $\Delta_k$  is swapped.
- DFP is not as popular as BFGS. There is some evidence that BFGS is more efficient than DFP.

# Outline

## 1 Optimization Problem

## 2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

## 3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and  
Davidon-Fletcher-Powell (DFP)

## 4 Stochastic Gradient Descent

## 5 Comparison Methods: Nelder-Mead Algorithm



# Steepest Ascent

The greatest possible increase in  $\bar{\ell}_T(\theta)$  for the (small enough) distance between  $\theta_k$  and  $\theta_{k+1}$  is provided by

$$\theta_{k+1} = \theta_k + \lambda_k g_k.$$

Motivated by the Lagrangian:

$$L = \underbrace{\bar{\ell}_T(\theta_k) + (\theta_{k+1} - \theta_k)g_k}_{\text{1st-order Taylor expansion of } \bar{\ell}_T(\theta_{k+1})} - \frac{1}{2\lambda_k} \underbrace{[(\theta_{k+1} - \theta_k)'(\theta_{k+1} - \theta_k) - d]}_{\text{distance from } \theta_k \text{ to } \theta_{k+1} \text{ being } \sqrt{d}}$$

Can pick  $\lambda_k$  that maximizes  $\bar{\ell}_T(\theta_k + \lambda_k g_k)$  (line search).

- “Steepest ascent” is only attained in a neighborhood of  $\theta_k$ . Usually converges more slowly than BHHH.
- For minimization problems, this is called the **gradient descent**.

# Stochastic Gradient Descent

- Historically gradient descent is not popular in nonlinear or nonconvex optimization problems because it gets stuck at local minima.
- For deep learning, computing the gradient can be very demanding. One way to be more efficient is the **stochastic gradient descent (SGD)**.
- Instead of evaluating the gradient of full sample, we **subsample**  $m \ll N$  observations with replacement and compute

$$\theta_{k+1} = \theta_k - \lambda_k g_k^*.$$

where  $g_k^*$  denotes the gradient of the subsample evaluated at  $\theta_k$ .

- Practitioners prefer  $m$  small ( $m = 1$ ): cheaper to compute and avoids overfitting (Goodfellow et al., 2016).

# Outline

## 1 Optimization Problem

## 2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

## 3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and  
Davidon-Fletcher-Powell (DFP)

## 4 Stochastic Gradient Descent

## 5 Comparison Methods: Nelder-Mead Algorithm

# Comparison Methods





- **Gradient-based** methods are susceptible to converge at a local maximum. Can use a variety of starting values to investigate the issue.
- An alternative is **comparison-based** methods: compute objective function at several points and pick the one yielding the optimum value.
- Comparison methods better behave with **non-smooth** objective functions. Stochastic comparison methods are more likely to find global optimum (in theory).

# Comparison Methods

## Nelder-Mead Algorithm

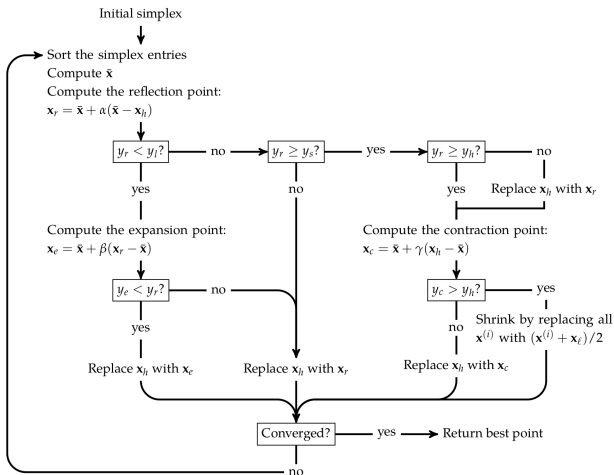
This is `fminsearch` in Matlab. Consider **minimizing** a generic criterion function  $y = f(x)$  with  $x \in \mathbb{R}^K$ .

- 1 Choose initial simplex  $\{x_1, x_2, \dots, x_{n+1}\}$ . Think of it as an  $n$ -dimensional version of a triangle.
- 2 Sort simplex vertices in descending order:  
 $f(x_h) > f(x_s) > \dots > f(x_l)$  (h: highest; s: second highest; l: lowest).
- 3 Modify the simplex in each step using one of the simplex operations.



# Comparison Methods

## Nelder-Mead Algorithm



# Bibliography

- Berndt, E. R., Hall, B. H., Hall, R. E., and Hausman, J. A. (1974), “Estimation and inference in nonlinear structural models,” in *Annals of Economic and Social Measurement, Volume 3, number 4*, NBER, pp. 653–665.
- Goodfellow, I., Bengio, Y., and Courville, A. (2016), *Deep Learning*, MIT Press, <http://www.deeplearningbook.org>.
- Kochenderfer, M. J. and Wheeler, T. A. (2019), *Algorithms for optimization*, Mit Press.
- Train, K. E. (2009), *Discrete Choice Methods with Simulation*, Cambridge University Press, 2nd ed.