

EC708 Discussion 8

Numerical Optimization

Yan Liu¹

Department of Economics
Boston University

March 18, 2022

¹Parts of the materials are from Train (2009), Kochenderfer and Wheeler (2019), and teaching slides of Shuowen Chen, Jean-Jacques Forneron, and Ryan Tibshirani.

Outline

1 Optimization Problem

2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and
Davidon-Fletcher-Powell (DFP)

4 Stochastic Gradient Descent

5 Comparison Methods: Nelder-Mead Algorithm

Optimization Problem

MLE

Consider maximizing the log-likelihood function


$$\bar{\ell}_T(\theta) = \sum_{t=1}^T \ell_t(\theta),$$

where $\theta \in \Theta \subset \mathbb{R}^K$.

- **Goal:** find $\hat{\theta}_T = \arg \max_{\theta \in \Theta} \bar{\ell}_T(\theta)$.
- To utilize minimization packages in practice, we usually work with $-\bar{\ell}_T(\theta)$.

Optimization Problem

Graphical Illustration of MLE



Finding $\hat{\theta}_T$ is a hill-climbing process:

- 1 Specify a starting value θ_0 .
- 2 Each step k moves to a new value θ_{k+1} at which $\bar{\ell}_T(\theta)$ is higher than at the current value θ_k .
- 3 Keep climbing until no further increase can be found.

Optimization Problem

Gradient and Hessian

Gradient in step k :

$$g_k = \left(\frac{\partial \bar{\ell}_T(\theta)}{\partial \theta} \right)_{\theta_k}.$$

Hessian in step k :

$$H_k = \left(\frac{\partial g_k}{\partial \theta'} \right)_{\theta_k} = \left(\frac{\partial^2 \bar{\ell}_T(\theta)}{\partial \theta \partial \theta'} \right)_{\theta_k}.$$

The gradient tells use **in what direction** to climb, and the Hessian can help us to know **how far** to climb.

Optimization Problem

Numerical Evaluation of Derivatives

Recall the definition of first-order partial derivative of function f at θ :

$$\frac{\partial f(\theta)}{\partial \theta_j} = \lim_{h \rightarrow 0} \frac{f(\theta_1, \dots, \theta_j + h, \dots, \theta_K) - f(\theta_1, \dots, \theta_j, \dots, \theta_K)}{h}$$

for $j = 1, \dots, K$. Numerically, we can approximate it by calculating

$$f_j(\theta) = \frac{f(\theta + he_j) - f(\theta)}{h},$$

where $e_j = (0, \dots, 0, 1, 0, \dots, 0)$ is the unit vector with 1 in position j and h is the step size. In practice, a more accurate approximation is

$$f_j(\theta) = \frac{f(\theta + he_j) - f(\theta - he_j)}{2h}.$$

Outline

1 Optimization Problem

2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and
Davidon-Fletcher-Powell (DFP)

4 Stochastic Gradient Descent

5 Comparison Methods: Nelder-Mead Algorithm

Newton-Raphson Method

Algorithm

Take a second-order Taylor's approximation of $\bar{\ell}_T(\theta_{k+1})$ around $\bar{\ell}_T(\theta_k)$:

$$\bar{\ell}_T(\theta_{k+1}) = \bar{\ell}_T(\theta_k) + (\theta_{k+1} - \theta_k)' g_k + \frac{1}{2} (\theta_{k+1} - \theta_k)' H_k (\theta_{k+1} - \theta_k).$$


Find θ_{k+1} that maximizes this approximation:

$$g_k + H_k(\theta_{k+1} - \theta_k) = 0 \Rightarrow \theta_{k+1} = \theta_k - \underbrace{H_k^{-1}}_{\text{step size}} \cdot \underbrace{g_k}_{\text{direction}}.$$

- Iterate until convergence, which can be defined in many ways:
 - $\bar{\ell}_T(\theta_{k+1})$ close to $\bar{\ell}_T(\theta_k)$
 - θ_{k+1} close to θ_k
 - g_{k+1} close to g_k
- If $\bar{\ell}_T(\theta)$ were **exactly quadratic**, then Newton-Raphson reaches the maximum in **one step** from any starting value.

Newton-Raphson Method

Step Size




- It is possible for Newton-Raphson to step past the maximum and move to a lower $\bar{\ell}_T(\theta)$.
- To ensure each step provides an increase in $\bar{\ell}_T(\theta)$, we introduce a scalar step size λ_k :

$$\theta_{k+1} = \theta_k - \lambda_k H_k^{-1} g_k.$$

Newton-Raphson Method

Determining the Step Size: Backtracking Line Search



Perform step-size adjustment in each step. Start with $\lambda_k = 1$.

- If $\bar{\ell}_T(\theta_{k+1}) < \bar{\ell}_T(\theta_k)$, continue halving λ_k until $\bar{\ell}_T(\theta_{k+1}) > \bar{\ell}_T(\theta_k)$.
 - A tiny λ_k is a signal that a different iteration procedure is needed.
- If $\bar{\ell}_T(\theta_{k+1}) > \bar{\ell}_T(\theta_k)$, continue doubling λ_k as long as doing so further raises $\bar{\ell}_T(\theta_{k+1})$.
 - Raising λ_k reduces # of iterations needed to reach the maximum.

Newton-Raphson Method

Drawbacks

- Calculations of the Hessian is usually computation-intensive.
 - $\frac{K(K+1)}{2}$ functions to evaluate in each step
 - Numerically calculated Hessian might be ill-behaved (singular).
- Does not guarantee an increase in each step if the log-likelihood function is not globally concave.
 - Hessian may not be negative definite.
 - Remedy: **regularization**. Instead of using H_k^{-1} directly, use

$$(H_k + \mu_k I_K)^{-1}$$

where $\mu_k < 0$ guarantees negative definiteness.

Gauss-Newton Method

Consider the following general nonlinear model:

$$y_t = f(x_t; \theta) + u_t, \quad t = 1, \dots, T$$

- $u_t \sim \text{i.i.d. } N(0, \sigma^2)$. Assume σ^2 is known here.
- x_t : $M \times 1$ exogenous regressors.
- $f(\cdot)$: some function satisfying some regularity conditions.

Gauss-Newton Method

Due to the normality assumption, we have the log-likelihood function

$$\bar{\ell}_T(\theta) = -\frac{T}{2} \ln 2\pi - \frac{T}{2} \ln \sigma^2 - \underbrace{\sum_{t=1}^T (y_t - f(x_t; \theta))^2}_{\equiv S(\theta)}.$$

It suffices to work with $S(\theta)$. Its first and second derivatives w.r.t. θ are

$$g(\theta) = 2 \sum_{t=1}^T \frac{\partial u_t}{\partial \theta} u_t, \quad H(\theta) = 2 \sum_{t=1}^T \left[\frac{\partial u_t}{\partial \theta} \frac{\partial u_t}{\partial \theta'} + \frac{\partial^2 u_t}{\partial \theta \partial \theta'} u_t \right].$$

In $H(\theta)$, the blue term is usually small relative to the red term, so we **neglect** it and use the following:

$$\theta_{k+1} = \theta_k - \left[\sum_{t=1}^T \frac{\partial u_t}{\partial \theta} \frac{\partial u_t}{\partial \theta'} \right]^{-1} \bigg|_{\theta=\theta_k} \sum_{t=1}^T \frac{\partial u_t}{\partial \theta} u_t \bigg|_{\theta=\theta_k}.$$

Gauss-Newton Method

Remarks:

- This method doesn't compute second-order derivatives.
- Has an OLS interpretation. Let $z_t = -\partial u_t / \partial \theta$, then

$$\theta_{k+1} = \theta_k + \left(\sum_{t=1}^T z_t z_t' \right)^{-1} \left| \sum_{t=1}^T z_t u_t \right|_{\theta=\theta_k}.$$

- Similar regularization can be incorporated as in Newton-Raphson:
Marquart quadratic hill climbing

$$\theta_{k+1} = \theta_k + \left(\sum_{t=1}^T z_t z_t' + \mu I_K \right)^{-1} \left| \sum_{t=1}^T z_t u_t \right|_{\theta=\theta_k}.$$

Outline

1 Optimization Problem

2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and Davidon-Fletcher-Powell (DFP)

4 Stochastic Gradient Descent

5 Comparison Methods: Nelder-Mead Algorithm

Quasi-Newton Methods

- Both main steps in Newton-Raphson method could be expensive:
 - Compute Hessian H_k
 - Solve the system $H_k \Delta = -g_k$
- Quasi-Newton methods repeat updates of the form

$$\theta_{k+1} = \theta_k - \lambda_k G_k^{-1} g_k$$

for some approximation G_k of H_k . We want G_k to be easy to compute and linear system $G_k \Delta = -g_k$ to be easy to solve.

Quasi-Newton Methods

Berndt-Hall-Hall-Hausman (BHHH)

Berndt et al. (1974) utilize the fact that the objective function is the sum of log likelihoods and propose to use scores to approximate Hessian.

- Score of observation t :

$$s_t(\theta_k) = \left. \frac{\partial \ell_t(\theta)}{\partial \theta} \right|_{\theta=\theta_k}.$$

- Gradient is the sum of scores:

$$g_k = \sum_{t=1}^T s_t(\theta_k).$$

Quasi-Newton Methods

Berndt-Hall-Hall-Hausman (BHHH)

- Outer product of observation t 's score is the $K \times K$ matrix

$$s_t(\theta_k)s_t(\theta_k)' = \begin{pmatrix} s_t^1 s_t^1 & s_t^1 s_t^2 & \cdots & s_t^1 s_t^K \\ s_t^2 s_t^1 & s_t^2 s_t^2 & \cdots & s_t^2 s_t^K \\ \vdots & \vdots & & \vdots \\ s_t^K s_t^1 & s_t^K s_t^2 & \cdots & s_t^K s_t^K \end{pmatrix},$$

where s_t^j is the j -th element of $s_t(\theta_k)$.

- Sum of outer product of scores:

$$G_k = \sum_{t=1}^T s_t(\theta_k)s_t(\theta_k)'.$$

Berndt-Hall-Hall-Hausman (BHHH) update uses G_k in place of $-H_k$:

$$\theta_{k+1} = \theta_k + \lambda_k G_k^{-1} g_k.$$

Quasi-Newton Methods

Berndt-Hall-Hall-Hausman (BHHH)

Why does BHHH work?

- At maximum, G_k is the sample variance of scores and thus provides a measure of the log-likelihood functions' curvature, similar to H_k .
- These ideas are formalized in the **information matrix equality**.
- G_k is far faster to calculate than H_k and necessarily positive definite.

Drawbacks: BHHH can give small steps when far from the maximum because G_k is not a good approximation to $-H_k$.

Quasi-Newton Methods

BFGS and DFP

- BHHH uses only information at θ_k to determine each step.
- As G_k already contains information about the Hessian, BFGS and DFP use suitable matrix update to form G_{k+1} .
- General procedure: In each iteration k ,
 - 1 Compute Quasi-Newton direction $\Delta_k = -G_k^{-1}g_k$
 - 2 Determine stepsize λ_k (by backtracking line search)
 - 3 Update $\theta_{k+1} = \theta_k + \lambda_k \Delta_k$
 - 4 Compute G_{k+1} from G_k

Quasi-Newton Methods

BFGS and DFP

Reasonable requirement for G_{k+1} (motivated by secant method):

$$g_{k+1} = g_k + G_{k+1} \Delta_k.$$

In addition, we want:

- G_{k+1} to be symmetric
- G_{k+1} to preserve positive definiteness (BFGS and DFP deal with convex optimization)

Quasi-Newton Methods

Broyden-Fletcher-Goldfarb-Shanno (BFGS)

Broyden-Fletcher-Goldfarb-Shanno (BFGS) update:

BFGS uses a rank-two update of the form:

$$G_{k+1} = G_k + auu' + bvv'.$$

Let $\gamma_k = g_{k+1} - g_k$. The secant equation yields

$$\gamma_k - G_k \Delta_k = (au' \Delta_k)u + (bv' \Delta_k)v.$$

Putting $u = \gamma_k$, $v = G_k \Delta_k$, and solving for a, b , we get

$$G_{k+1} = G_k - \frac{G_k \Delta_k \Delta_k' G_k}{\Delta_k' G_k \Delta_k} + \frac{\gamma_k \gamma_k'}{\gamma_k' \Delta_k}.$$

- BFGS update is quite cheap: $O(K^2)$ operations
- BFGS is the algorithm behind Matlab's **fminunc**.

Quasi-Newton Methods

Davidon-Fletcher-Powell (DFP)

Davidon-Fletcher-Powell (DFP) update:

DFP pursues the same idea to update G_{k+1}^{-1} :

$$G_{k+1}^{-1} = G_k^{-1} + auu' + bvv'.$$

The secant equation yields

$$\Delta_k - G_k^{-1}\gamma_k = (au'\gamma_k)u + (bv'\gamma_k)v.$$

Putting $u = \Delta_k$, $v = G_k^{-1}\gamma_k$, and solving for a, b , we get

$$G_{k+1}^{-1} = G_k^{-1} - \frac{G_k^{-1}\gamma_k\gamma_k'G_k^{-1}}{\gamma_k'G_k^{-1}\gamma_k} + \frac{\Delta_k\Delta_k'}{\Delta_k'\gamma_k}.$$

- The role of γ_k and Δ_k is swapped.
- DFP is not as popular as BFGS. There is some evidence that BFGS is more efficient than DFP.

Outline

1 Optimization Problem

2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and
Davidon-Fletcher-Powell (DFP)

4 Stochastic Gradient Descent

5 Comparison Methods: Nelder-Mead Algorithm

Steepest Ascent

The greatest possible increase in $\bar{\ell}_T(\theta)$ for the (small enough) distance between θ_k and θ_{k+1} is provided by

$$\theta_{k+1} = \theta_k + \lambda_k g_k.$$

Motivated by the Lagrangian:

$$L = \underbrace{\bar{\ell}_T(\theta_k) + (\theta_{k+1} - \theta_k)g_k}_{\text{1st-order Taylor expansion of } \bar{\ell}_T(\theta_{k+1})} - \frac{1}{2\lambda_k} \underbrace{[(\theta_{k+1} - \theta_k)'(\theta_{k+1} - \theta_k) - d]}_{\text{distance from } \theta_k \text{ to } \theta_{k+1} \text{ being } \sqrt{d}}$$

Can pick λ_k that maximizes $\bar{\ell}_T(\theta_k + \lambda_k g_k)$ (line search).

- “Steepest ascent” is only attained in a neighborhood of θ_k . Usually converges more slowly than BHHH.
- For minimization problems, this is called the **gradient descent**.

Stochastic Gradient Descent

- Historically gradient descent is not popular in nonlinear or nonconvex optimization problems because it gets stuck at local minima.
- For deep learning, computing the gradient can be very demanding. One way to be more efficient is the **stochastic gradient descent (SGD)**.
- Instead of evaluating the gradient of full sample, we **subsample** $m \ll N$ observations with replacement and compute

$$\theta_{k+1} = \theta_k - \lambda_k g_k^*.$$

where g_k^* denotes the gradient of the subsample evaluated at θ_k .

- Practitioners prefer m small ($m = 1$): cheaper to compute and avoids overfitting (Goodfellow et al., 2016).

Outline

1 Optimization Problem

2 Full-Newton Method

- Newton-Raphson
- Gauss-Newton

3 Quasi-Newton Methods

- Berndt-Hall-Hall-Hausman (BHHH)
- Broyden-Fletcher-Goldfarb-Shanno (BFGS) and
Davidon-Fletcher-Powell (DFP)

4 Stochastic Gradient Descent

5 Comparison Methods: Nelder-Mead Algorithm

Comparison Methods





- **Gradient-based** methods are susceptible to converge at a local maximum. Can use a variety of starting values to investigate the issue.
- An alternative is **comparison-based** methods: compute objective function at several points and pick the one yielding the optimum value.
- Comparison methods better behave with **non-smooth** objective functions. Stochastic comparison methods are more likely to find global optimum (in theory).

Comparison Methods

Nelder-Mead Algorithm

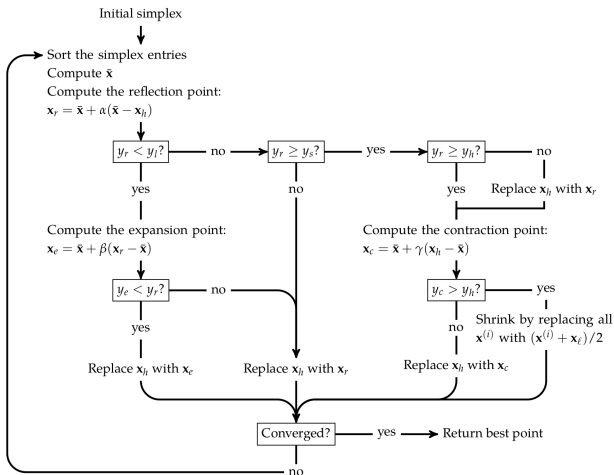
This is `fminsearch` in Matlab. Consider **minimizing** a generic criterion function $y = f(x)$ with $x \in \mathbb{R}^K$.

- 1 Choose initial simplex $\{x_1, x_2, \dots, x_{n+1}\}$. Think of it as an n -dimensional version of a triangle.
- 2 Sort simplex vertices in descending order:
 $f(x_h) > f(x_s) > \dots > f(x_l)$ (h: highest; s: second highest; l: lowest).
- 3 Modify the simplex in each step using one of the simplex operations.



Comparison Methods

Nelder-Mead Algorithm



Bibliography

- Berndt, E. R., Hall, B. H., Hall, R. E., and Hausman, J. A. (1974), “Estimation and inference in nonlinear structural models,” in *Annals of Economic and Social Measurement, Volume 3, number 4*, NBER, pp. 653–665.
- Goodfellow, I., Bengio, Y., and Courville, A. (2016), *Deep Learning*, MIT Press, <http://www.deeplearningbook.org>.
- Kochenderfer, M. J. and Wheeler, T. A. (2019), *Algorithms for optimization*, Mit Press.
- Train, K. E. (2009), *Discrete Choice Methods with Simulation*, Cambridge University Press, 2nd ed.